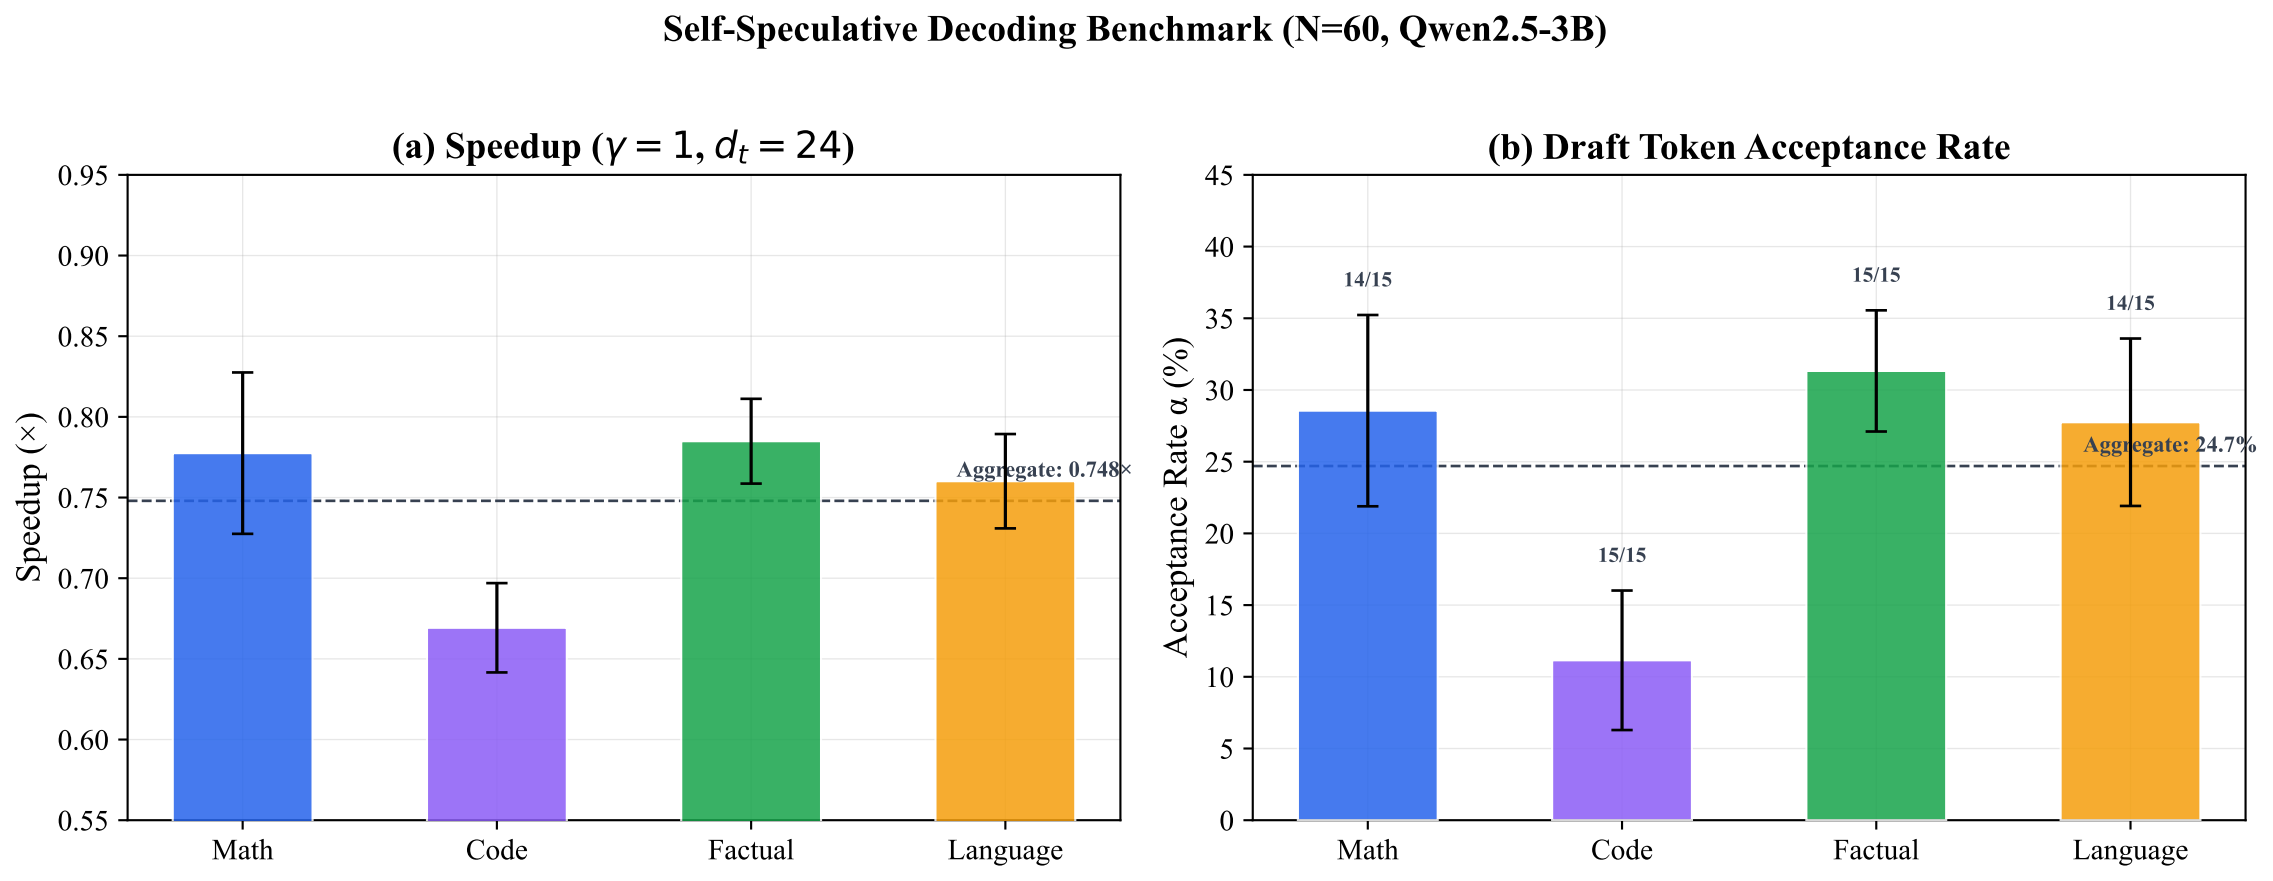
Self-Speculative Decoding Benchmark (N=60, Qwen2.5-3B)
(a) Speedup ($\gamma = 1$, $d_t = 24$)
Speedup (×)
0.95
0.90
0.85
0.80
0.75
0.70
0.65
0.60
0.55
Aggregate: 0.748×
Math
Code
Factual
Language
(b) Draft Token Acceptance Rate
Acceptance Rate α (%)
45
40
35
30
25
20
15
10
5
0
14/15
15/15
15/15
14/15
Aggregate: 24.7%
Math
Code
Factual
Language